

Bielefeld University

Faculty of Technology

Semantic Computing Group

Thesis Outline

*Neuro-Symbolic Framework for
Common-Sense Reasoning in Domestic Robots*

Author: Cono Cirone

Supervisor: Prof. Dr. Philipp Cimiano

Advisor: Jan-Philipp Töberg

June 2026

Contents

1	Introduction	4
1.1	Context and Motivation: Autonomous Agents in Domestic Environments	4
1.2	The Challenge of Embodied Common-Sense Knowledge	4
1.3	Limitations of Pure LLM Approaches	4
1.4	Research Questions and Objectives	4
1.5	Contributions	4
2	Background and Related Work	4
2.1	Common-Sense Knowledge in Robotics	4
2.1.1	Definition and Scope of Embodied Common-Sense	4
2.1.2	Knowledge Graphs: ConceptNet	4
2.1.3	Knowledge Graphs: CSKG	4
2.1.4	Ontologies for the Domestic Domain (SOMA)	4
2.2	Large Language Models and Reasoning	4
2.2.1	LLMs as Knowledge Sources	4
2.3	Neuro-Symbolic AI	4
2.3.1	Combining Neural and Symbolic Approaches	4
2.3.2	Inductive Logic Programming (ILASP)	4
2.3.3	Logic Programming for Runtime Reasoning (SWI-Prolog, Clingo/ASP)	4
3	The RoboCSKBench Benchmark	4
3.1	Benchmark Design and Rationale	4
3.2	Task Descriptions	4
3.2.1	Tidy Up—Object Placement (Generative and Multiple Choice)	4
3.2.2	Tool Usage—Affordance-Based Tool Selection	4
3.2.3	Meta-Reasoning—Robot Hardware Capability Assessment . . .	4
3.2.4	Table Setting—Cutlery and Plate Prediction	4
3.2.5	Procedural Knowledge—Temporal Step Ordering	4
3.3	Data Sources and Construction	4
3.4	Evaluation Metrics (Accuracy, MRR, F1, Jaccard Similarity)	4
4	Pipeline 1: LLM → ILASP (Offline Rule Learning)	5
4.1	Architecture Overview	5
4.2	Knowledge Base Construction (Offline Phase)	5
4.2.1	ConceptNet Subgraph Extraction	5
4.2.2	CSKG Subgraph Extraction	5
4.2.3	Semantic Object Filtering via LLM	5
4.2.4	SOMA Ontology Annotation via LLM	5
4.2.5	ASP Integrity Constraints for Validation	5
4.3	Inductive Rule Learning	5
4.3.1	ILASP Example Generation	5
4.3.2	LLM-Guided Rule Induction with Clingo Verification	5
4.3.3	Rule Pruning and Minimalization	5
4.4	Online Inference and Evaluation	5
4.4.1	Clingo-Based Inference on Unseen Objects	5
4.4.2	Results on Tidy Up (Robo-CSK-Bench)	5

4.4.3	Results on Tool Usage (Robo-CSK-Bench)	5
4.4.4	Discussion	5
5	Pipeline 2: LLM $\rightarrow$ SWI-Prolog (Runtime Neuro-Symbolic Reasoning)	5
5.1	Architecture Overview: Neural $\rightarrow$ Symbolic $\rightarrow$ Neural	5
5.2	The General Architecture	5
5.2.1	Neural Phase: LLM Extracts Semantic Properties	5
5.2.2	Symbolic Phase: Prolog Applies Domain Rules	5
5.2.3	Neural Phase: LLM Makes the Final KB-Grounded Decision	5
5.3	Task-Specific Implementations	5
5.3.1	Tidy Up — Location KB + Prolog Location Selector	5
5.3.2	Tool Usage — Affordance KB + Prolog Tool Selector	5
5.3.3	Meta-Reasoning — Robot Capability Axioms in Prolog	5
5.3.4	Table Setting — Meal Property Extraction + Prolog Cutlery/Plate Rules	5
5.3.5	Procedural Knowledge — Temporal Phase Reasoner in Prolog	5
6	Experimental Evaluation	6
6.1	Experimental Setup	6
6.1.1	Models and Configurations	6
6.1.2	Ablation Study Design (Full Pipeline, Pure LLM, Pure Logic, No CoT)	6
6.2	Results on Robo-CSK-Bench (In-Distribution)	6
6.2.1	Tidy Up — Multiple Choice and Open/Generative	6
6.2.2	Tool Usage	6
6.2.3	Meta-Reasoning — Binary and Multiple Choice	6
6.2.4	Table Setting — Plate Accuracy and Cutlery Jaccard	6
6.2.5	Procedural Knowledge — Binary and Multiple Choice	6
6.3	Ablation Analysis	6
6.3.1	Contribution of the Symbolic Layer	6
6.3.2	Neuro-Symbolic vs. Pure LLM Delta	6
6.3.3	When Does the Symbolic Layer Help Most?	6
6.4	Comparison with Benchmark Paper Baselines (GPT-4o, Llama 3.3, Gemma 2)	6
6.5	Cross-Task Summary and Discussion	6
7	Generalization Study (Out-of-Distribution Evaluation)	7
7.1	Motivation and Study Design	7
7.2	OOD Datasets	7
7.2.1	VirtualHome and TidyBot (Tidy Up)	7
7.2.2	PIQA — Household-Filtered (Tool Usage)	7
7.2.3	BEHAVIOR-1K (Meta-Reasoning)	7
7.2.4	RecipeNLG (Procedural Knowledge)	7
7.3	Results	7
7.3.1	Tidy Up Generalization (VirtualHome and TidyBot)	7
7.3.2	Tool Usage Generalization (PIQA)	7
7.3.3	Meta-Reasoning Generalization (BEHAVIOR-1K)	7
7.3.4	Procedural Knowledge Generalization (RecipeNLG)	7
7.4	OOD Performance Summary and Discussion	7

7.4.1	Where the Full Pipeline Excels on OOD Data	7
7.4.2	Where Pure LLM Remains Competitive	7
7.4.3	Zero-Shot Transfer of Symbolic Rules	7
8	Conclusions	7
8.1	Summary of Contributions	7
8.2	Key Findings	7
8.3	Limitations	7
8.4	Future Work	7
	References	7
	Appendices	7

1. Introduction

- 1.1 Context and Motivation: Autonomous Agents in Domestic Environments
- 1.2 The Challenge of Embodied Common-Sense Knowledge
- 1.3 Limitations of Pure LLM Approaches
- 1.4 Research Questions and Objectives
- 1.5 Contributions

2. Background and Related Work

- 2.1 Common-Sense Knowledge in Robotics
 - 2.1.1 Definition and Scope of Embodied Common-Sense
 - 2.1.2 Knowledge Graphs: ConceptNet
 - 2.1.3 Knowledge Graphs: CSKG
 - 2.1.4 Ontologies for the Domestic Domain (SOMA)
- 2.2 Large Language Models and Reasoning
 - 2.2.1 LLMs as Knowledge Sources
- 2.3 Neuro-Symbolic AI
 - 2.3.1 Combining Neural and Symbolic Approaches
 - 2.3.2 Inductive Logic Programming (ILASP)
 - 2.3.3 Logic Programming for Runtime Reasoning (SWI-Prolog, Clingo/ASP)

3. The RoboCSKBench Benchmark

- 3.1 Benchmark Design and Rationale
- 3.2 Task Descriptions
 - 3.2.1 Tidy Up—Object Placement (Generative and Multiple Choice)
 - 3.2.2 Tool Usage—Affordance-Based Tool Selection
 - 3.2.3 Meta-Reasoning—Robot Hardware Capability Assessment
 - 3.2.4 Table Setting—Cutlery and Plate Prediction
 - 3.2.5 Procedural Knowledge—Temporal Step Ordering
- 3.3 Data Sources and Construction
- 3.4 Evaluation Metrics (Accuracy, MRR, F1, Jaccard Similarity)

4. Pipeline 1: LLM $\rightarrow$ ILASP (Offline Rule Learning)

4.1 Architecture Overview

4.2 Knowledge Base Construction (Offline Phase)

4.2.1 ConceptNet Subgraph Extraction

4.2.2 CSKG Subgraph Extraction

4.2.3 Semantic Object Filtering via LLM

4.2.4 SOMA Ontology Annotation via LLM

4.2.5 ASP Integrity Constraints for Validation

4.3 Inductive Rule Learning

4.3.1 ILASP Example Generation

4.3.2 LLM-Guided Rule Induction with Clingo Verification

4.3.3 Rule Pruning and Minimalization

4.4 Online Inference and Evaluation

4.4.1 Clingo-Based Inference on Unseen Objects

4.4.2 Results on Tidy Up (Robo-CSK-Bench)

4.4.3 Results on Tool Usage (Robo-CSK-Bench)

4.4.4 Discussion

5. Pipeline 2: LLM $\rightarrow$ SWI-Prolog (Runtime Neuro-Symbolic Reasoning)

5.1 Architecture Overview: Neural $\rightarrow$ Symbolic $\rightarrow$ Neural

5.2 The General Architecture

5.2.1 Neural Phase: LLM Extracts Semantic Properties

5.2.2 Symbolic Phase: Prolog Applies Domain Rules

5.2.3 Neural Phase: LLM Makes the Final KB-Grounded Decision

5.3 Task-Specific Implementations

5.3.1 Tidy Up — Location KB + Prolog Location Selector

5.3.2 Tool Usage — Affordance KB + Prolog Tool Selector

5.3.3 Meta-Reasoning — Robot Capability Axioms in Prolog

5.3.4 Table Setting — Meal Property Extraction + Prolog Cutlery/Plate Rules

5.3.5 Procedural Knowledge — Temporal Phase Reasoner in Prolog

6. Experimental Evaluation

6.1 Experimental Setup

6.1.1 Models and Configurations

6.1.2 Ablation Study Design (Full Pipeline, Pure LLM, Pure Logic, No CoT)

6.2 Results on Robo-CSK-Bench (In-Distribution)

6.2.1 Tidy Up—Multiple Choice and Open/Generative

6.2.2 Tool Usage

6.2.3 Meta-Reasoning—Binary and Multiple Choice

6.2.4 Table Setting—Plate Accuracy and Cutlery Jaccard

6.2.5 Procedural Knowledge—Binary and Multiple Choice

6.3 Ablation Analysis

6.3.1 Contribution of the Symbolic Layer

6.3.2 Neuro-Symbolic vs. Pure LLM Delta

6.3.3 When Does the Symbolic Layer Help Most?

6.4 Comparison with Benchmark Paper Baselines (GPT-4o, Llama 3.3, Gemma 2)

6.5 Cross-Task Summary and Discussion

7. Generalization Study (Out-of-Distribution Evaluation)

7.1 Motivation and Study Design

7.2 OOD Datasets

7.2.1 VirtualHome and TidyBot (Tidy Up)

7.2.2 PIQA — Household-Filtered (Tool Usage)

7.2.3 BEHAVIOR-1K (Meta-Reasoning)

7.2.4 RecipeNLG (Procedural Knowledge)

7.3 Results

7.3.1 Tidy Up Generalization (VirtualHome and TidyBot)

7.3.2 Tool Usage Generalization (PIQA)

7.3.3 Meta-Reasoning Generalization (BEHAVIOR-1K)

7.3.4 Procedural Knowledge Generalization (RecipeNLG)

7.4 OOD Performance Summary and Discussion

7.4.1 Where the Full Pipeline Excels on OOD Data

7.4.2 Where Pure LLM Remains Competitive

7.4.3 Zero-Shot Transfer of Symbolic Rules

8. Conclusions

8.1 Summary of Contributions

8.2 Key Findings

8.3 Limitations

8.4 Future Work

References

Appendices

A. SOMA Ontology Codebook and Annotation Schema

B. Prolog Reasoner Specifications (per Task)

C. Full Results Tables

D. Knowledge Base Statistics and Coverage